

**VGG trained on ImageNet was used to extract features and transfer learning was performed**

**The resulting features were 10 x 10 x 512. i.e. a total of 51,200 when flattened. These features were fed to a regular dense ANN producing Mean Sq.Error of 0.07 which was much better than our CNN's M.S.E. of 0.24. Prima facie this seems great, VGG has been trained on millions of images and can obviously detect edges, lines and all sorts of geometric properties better than the convolution layers in our custom CNN.**

**However our CNN produced MSE of 0.21 on test set whereas transfer learnt model produced 0.19. This indicates VGG transfer Learnt Model severely overfits.**

	Training Mean Squared Error	Test Mean Squared Error	
<b>Our CNN</b>	0.24	0.21	
<b>VGG Transfer Learnt</b>	0.07	0.19	Overfitting

```
In [4]: import numpy as np
import PIL, os
from PIL import Image
import glob, random
import tensorflow as tf
from tensorflow.keras import datasets, layers, models
import matplotlib.pyplot as plt
from sklearn import metrics

import seaborn as sns

from keras.applications.vgg16 import VGG16
```

```
In [41]: #np.savez_compressed('traintest.npz', X_train=X_train, y_train=y_train, X_test=X_test, y_test=y_test, names_train = names_train, names_test = names_test)
```

```
In [6]: traintest = np.load('traintest.npz')
```

```
In [7]: traintest.files
```

```
Out[7]: ['X_train', 'y_train', 'X_test', 'y_test', 'names_train', 'names_test']
```

```
In [8]: traintest['names_test']
```

```
Out[8]: array(['AF1.jpg', 'AF100.jpg', 'AF1007.jpg', ..., 'CM84.jpg', 'CM88.jpg',
               'CM96.jpg'], dtype='<U10')
```

```
In [9]: X_train = traintest['X_train']
X_test = traintest['X_test']
y_train = traintest['y_train']
y_test = traintest['y_test']
```

```
In [18]: X_train.shape
```

```
Out[18]: (4400, 350, 350, 3)
```

## Feature Extractor

```
In [14]: #####
#Load model without classifier/fully connected layers
VGG_model = VGG16(weights='imagenet', include_top=False, input_shape=(350, 350, 3))

#Make loaded layers as non-trainable. This is important as we want to work with pre-trained weights
for layer in VGG_model.layers:
    layer.trainable = False
VGG_model.summary() #Trainable parameters will be 0
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_4 (InputLayer)	[(None, 350, 350, 3)]	0
block1_conv1 (Conv2D)	(None, 350, 350, 64)	1792
block1_conv2 (Conv2D)	(None, 350, 350, 64)	36928
block1_pool (MaxPooling2D)	(None, 175, 175, 64)	0
block2_conv1 (Conv2D)	(None, 175, 175, 128)	73856
block2_conv2 (Conv2D)	(None, 175, 175, 128)	147584
block2_pool (MaxPooling2D)	(None, 87, 87, 128)	0
block3_conv1 (Conv2D)	(None, 87, 87, 256)	295168
block3_conv2 (Conv2D)	(None, 87, 87, 256)	590080
block3_conv3 (Conv2D)	(None, 87, 87, 256)	590080
block3_pool (MaxPooling2D)	(None, 43, 43, 256)	0
block4_conv1 (Conv2D)	(None, 43, 43, 512)	1180160
block4_conv2 (Conv2D)	(None, 43, 43, 512)	2359808
block4_conv3 (Conv2D)	(None, 43, 43, 512)	2359808
block4_pool (MaxPooling2D)	(None, 21, 21, 512)	0
block5_conv1 (Conv2D)	(None, 21, 21, 512)	2359808
block5_conv2 (Conv2D)	(None, 21, 21, 512)	2359808
block5_conv3 (Conv2D)	(None, 21, 21, 512)	2359808

---

```

block5_pool (MaxPooling2D)      (None, 10, 10, 512)      0
=====
Total params: 14,714,688
Trainable params: 0
Non-trainable params: 14,714,688

```

---

```

In [20]: #Now, let us use features from convolutional network for random forest or ANN
feature_extractor=VGG_model.predict(X_train, verbose=1)

```

```

138/138 [=====] - 2340s 17s/step

```

```

In [21]: ann = models.Sequential([
    # each filter is a feature detector. Convolving feature detector with original image gives feature map
    # layers.Conv2D(filters=100, kernel_size=(10, 10), activation='relu', input_shape=(350, 350, 3)),
    # layers.MaxPooling2D((10, 10)),

    # layers.Conv2D(filters=50, kernel_size=(10, 10), activation='relu'),
    # layers.MaxPooling2D((10, 10)),

    # conventional ann
    layers.Flatten(input_shape=(10, 10, 512)),
    layers.Dense(30, activation='relu'),
    layers.Dense(10, activation='relu'),
    layers.Dense(1)
])

```

```

In [22]: ann.compile(optimizer='adam',
                    loss="mean_squared_error",
                    metrics=["mean_squared_error"])

```

```

In [ ]: features = feature_extractor.reshape(feature_extractor.shape[0], -1)

X_for_RF = features #This is our X input to RF

```

```
In [23]: ann.fit(feature_extractor, y_train, epochs=5)
```

```
Epoch 1/5  
138/138 [=====] - 3s 8ms/step - loss: 2.1897 - mean_squared_error: 2.1897  
Epoch 2/5  
138/138 [=====] - 1s 7ms/step - loss: 0.1551 - mean_squared_error: 0.1551  
Epoch 3/5  
138/138 [=====] - 1s 7ms/step - loss: 0.1162 - mean_squared_error: 0.1162  
Epoch 4/5  
138/138 [=====] - 1s 7ms/step - loss: 0.0911 - mean_squared_error: 0.0911  
Epoch 5/5  
138/138 [=====] - 1s 7ms/step - loss: 0.0782 - mean_squared_error: 0.0782
```

```
Out[23]: <keras.callbacks.History at 0x56bd4c520>
```

```
In [24]: # 1st do feature transformations on X_test
```

```
feature_extractor_test=VGG_model.predict(X_test, verbose=1)  
  
35/35 [=====] - 572s 16s/step
```

```
In [26]: ann.evaluate(feature_extractor_test,y_test)
```

```
35/35 [=====] - 0s 3ms/step - loss: 0.1906 - mean_squared_error: 0.1906
```

```
Out[26]: [0.19063855707645416, 0.19063855707645416]
```

-----END-----

